

⚠ Please do not reach out to reviewers directly; messaging them will not speed up your review.

Long Context Task Checklist

Use this checklist before submitting or accepting any task labeled

`long_context`.

Spec Anchor

A `long_context` task must require the model to test its context window by reading large documents. The file/corpus must be at least 50k tokens and cannot simply be parsed programmatically or solved through keyword search.

Allowed file/content examples include markdown, txt, PDF, docx, HTML, notebooks, chat logs, email threads, academic papers, documentation, and logs. The file extension is not enough; the content must actually function as long document context.

Required Checks

A task should only be labeled `long_context` if all of these are true:

1. It contains at least 50k tokens of valid document-like long-context content.
2. The long-context files are shipped with the task, not only generated by scripts during setup.
3. The long documents are authoritative to the solution.
4. The agent must read and reason over the documents to solve the task correctly.
5. The task cannot be solved by simple keyword search, grep, field extraction, or top-k statistics.
6. The long content is not primarily JSON, JSONL, CSV, TSV, database dumps, or uniform table/log records.

7. The corpus is not filler, repeated boilerplate, random text, or many tiny one-line docs.
8. The verifier checks behavior or outputs that depend on details from the long documents.
9. If multiple documents are provided, the task requires resolving interactions across documents, not just opening one obvious file.
10. The instruction tells the agent where the long documents live without leaking the exact answer path.

Immediate Rejection Patterns

Reject or relabel the task if it has any of these patterns:

- Under 50k valid long-context tokens.
- A generator script but no shipped document corpus.
- Large JSONL/CSV/TSV/database inputs labeled as long context.
- Uniform machine logs with clear delimiters and fixed fields.
- Dozens of tiny files whose aggregate content is far below 50k tokens.
- Long files used only for keyword search or simple analytics.
- Long documents that are historical, optional, non-authoritative, or unrelated to the oracle.
- Repeated/filler text used to inflate size.
- The main challenge is writing a parser, scheduler, ETL job, search endpoint, or analytics pipeline over structured data.

Good Task Shape

A strong `long_context` task usually looks like one of these:

- A long authoritative handbook/manual/spec that defines many interacting implementation rules.
- A long policy or underwriting manual where rates, exceptions, and edge cases must be resolved from text.
- A multi-file investigation corpus with emails, chats, calendars, HR records, and notes that must be reconciled.
- A long incident narrative/log corpus where the challenge is reconstructing what happened, not parsing a fixed schema.

EC Self-Review Questions

Before submitting, the EC should be able to answer yes to all of these:

- Could an agent solve this without reading the long documents? If yes, it is not long-context.
- Are the documents the source of truth rather than supporting decoration? If no, revise.
- Is the corpus over 50k valid document tokens after excluding code/config/structured data? If no, revise.
- Is the solution more than keyword search or simple parsing? If no, relabel or redesign.
- Would a reviewer understand why this tests long-context reading from the instruction and files? If no, revise.

PREVIOUS

[Reviewer Checklist](#)

NEXT

[Common Errors](#)